



Automate Everything

Hardware as Code

Hi I'm Tim



- “Docker Captain” – Community Leader
- PHP JHB Organiser
- 17 Years Development and DevOps
- tim@haak.co

Quick Survey

- Primary laptop/desktop OS?
- Primary server OS?
- Who has set up a server?
- Who finds setting up servers fun?



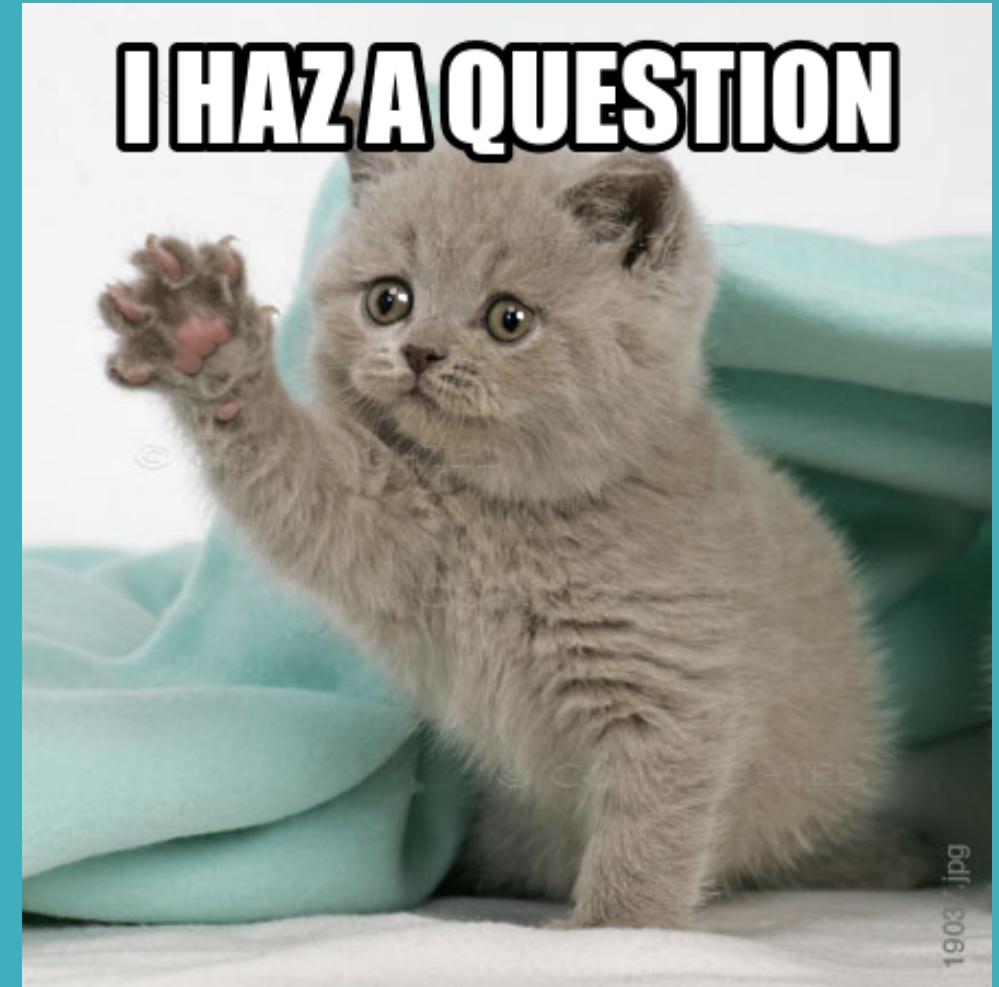
19037.jpg

First, what are the problems?



How long...

- Provision a new server?
- Rebuild all the pieces for a service?
- Work out what is on a server?
- Rollout updates to everything?
- Roll back a service or server?



Problem 1: Speed

- Setting up servers by hand can take quite a while.
- Deploying new services
 - Slow by hand
 - Error prone
- Takes too long to set up a clean server to quickly test something



Problem 2: What's really on the server?

- Human memory problems
 - What was on this server again?
- Is live, staging, and testing the same?
- But it worked on the developers/my PC!

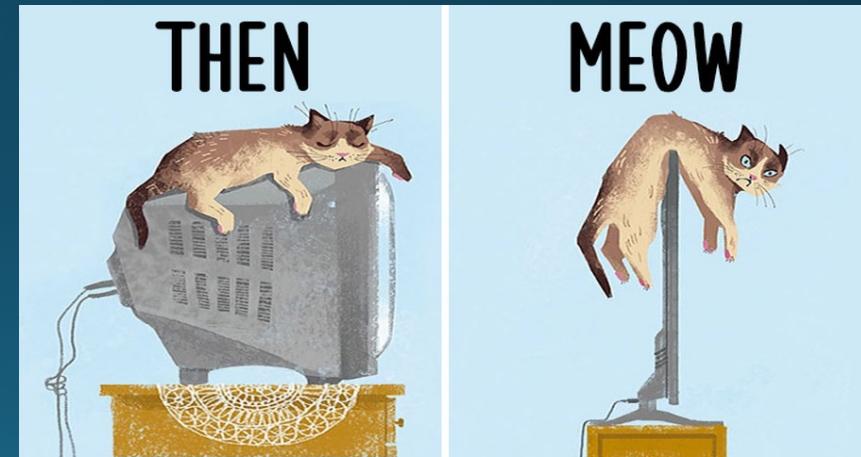


Problem 2: What's really on the server?

- Bob's sick, how did he set this up?
 - How do I bring new people on? (Documentation)
- We are error prone
 - PC's only ever do what we tell them.

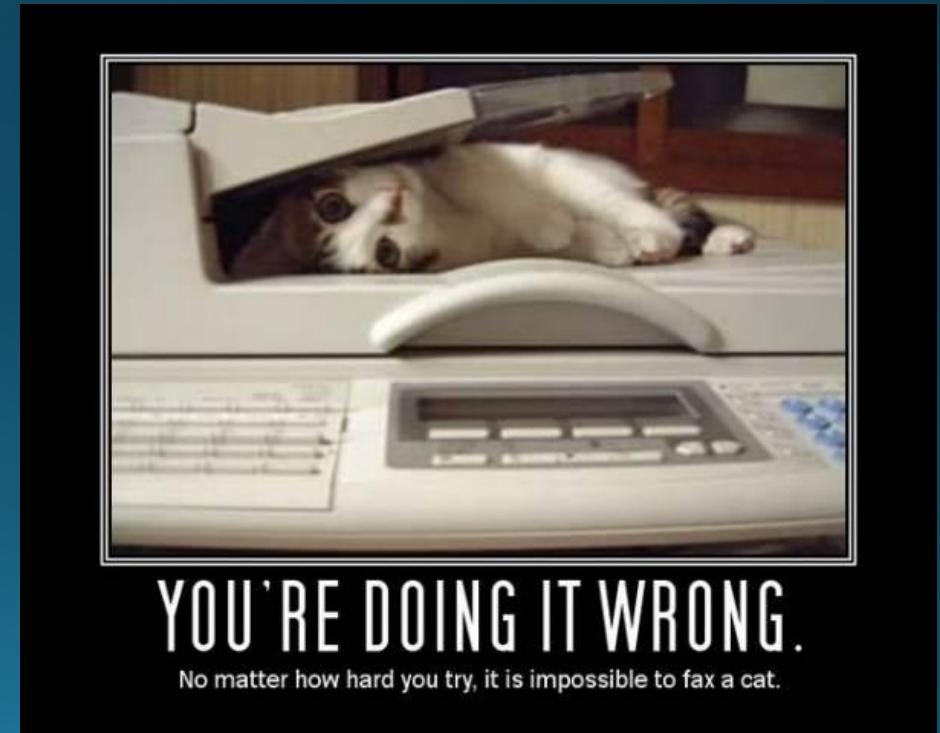
Problem 3: What Changed?

- What was changed?
- When was it changed?
- Can I prove either of these?



Problem 4: Repeatability

- I'm not always sure when I roll out that it will work
 - So let's not do it often



Bad patterns

- Cowboy/Cowgirl style
 - Just do it! We'll worry when it breaks...
 - I don't test, but when I do I do it on live!



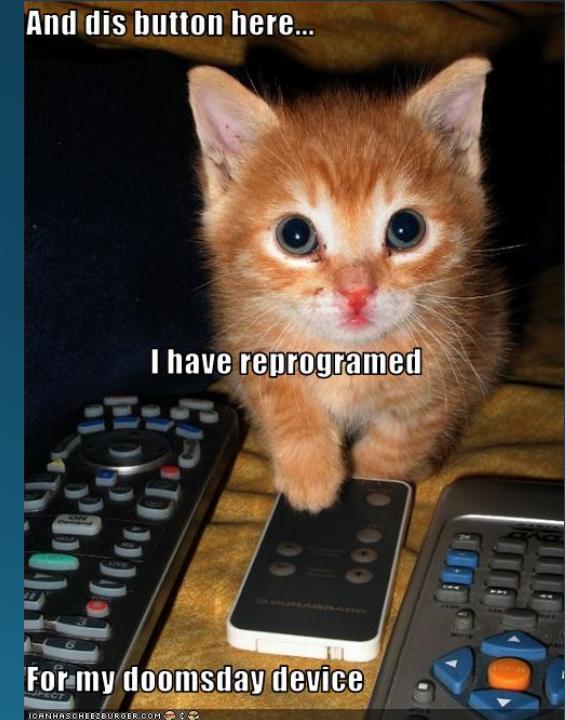
Bad patterns

- Bureaucrat style
 - If it's breaking, it means we don't have enough documentation/processes.
 - We can't roll out the critical patch to fix the broken thing as change control only meets next week. (True story)



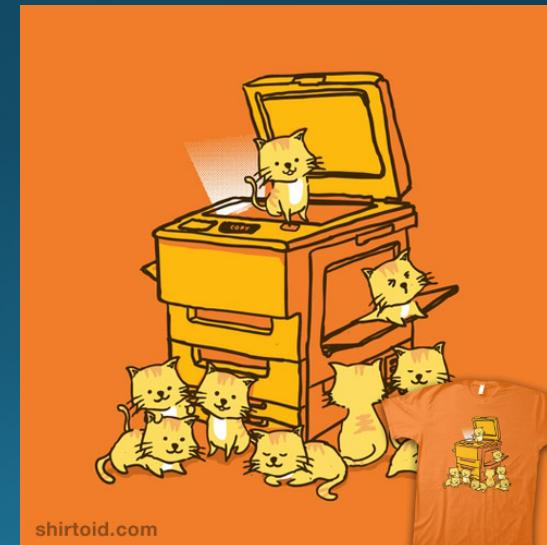
Better patterns to solve these

- Automation
 - Automate all the things!!!
 - Have tools that everyone in the team can use
 - But you know are safe and tested.
 - Have an option to easily roll back.
 - Blue/Green deployments.
 - Add web based frontends to these
 - DevOps don't just tell people to run a script.
 - Dev you have to help write the frontend to run the script.

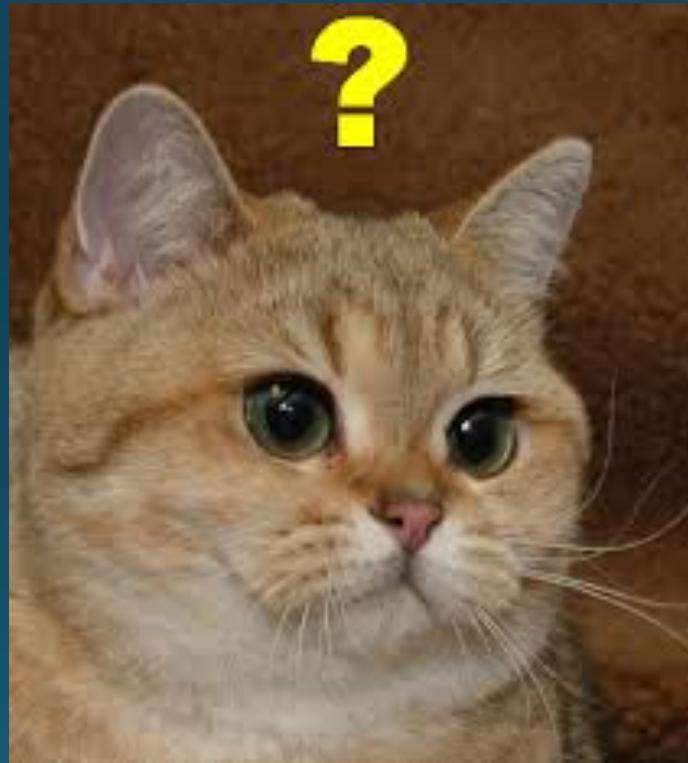


Better patterns to solve these

- Immutability where we can
 - An **immutable object** is an object whose state cannot be modified after it is created



But how do we do this?



But how do we do this?

- On the server, hardware and operating system side
 - **Hardware/Infrastructure as code**
 - Describe layout in text files that are committed to version control systems
 - Continuous Integration (CI)
 - Automated testing
 - Automated building

But how do we do this?

- On the code side
 - Code is text that you commit to version control systems
 - Continuous Integration (CI)
 - Automated testing
 - Automated building

These look strangely similar

- On the server, hardware and operating system side
 - **Hardware/Infrastructure as code**
 - Describe layout in text files that are committed to version control systems
 - Continuous Integration (CI)
 - Automated testing
 - Automated building
- On the code side
 - Code is text that you commit to version control systems
 - Continuous Integration (CI)
 - Automated testing
 - Automated building

How do we do this practically?



Step 1:Definition Files

Use Definition Files

- All configuration is defined in executable configuration definition files, such as shell scripts, [Ansible](#), [SaltStack](#), [Chef](#), or [Puppet](#).
- Personal preference is [Ansible](#) as it's simpler to get started.
- Commit scripts to your version control.

Cattle not pets

- Servers should be treated as cattle not pets!



Stop doing things by hand

- At no time should anyone log into a server and make on-the-fly adjustments.
- Any time you do this you risk turning your server into a pet.



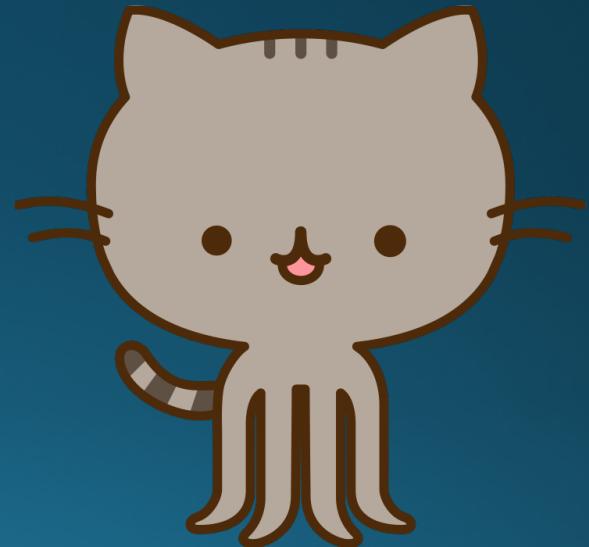
Bonus outcome: Documentation

- You've just documented what is on your server.
- If you need it in a more human format it can be generated from the code.
- Unlike a document a human would follow and run by hand, you can guarantee that it was done exactly as set in the code.



Version All The Things!!!

- Keep this code in version control.
- Changes can now be audited.
- Something breaking on live?
 - Just use the script to build a test server.



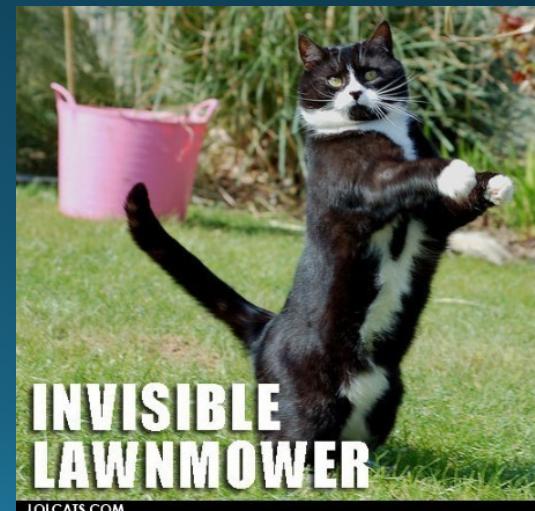
Continuous Deployment and Testing

- Now that you have it as code you can use all the normal code tooling to run and test.

Run It Often

- The smaller the change the lower the risk.
 - Fewer things to keep in your head.
- Ensure that your servers are always in a know state (Immutable).
- The more often you do something...The more predictable it gets.
- Doing things more often actually decreases risk.

Step 2: Build The Whole Server



INVISIBLE
LAWN MOWER

LOLCATS.COM

Use Definition Files for Everything!

- Stop doing the initial operating system install by hand.
- Here I use Packer.
 - Handles most cloud providers and virtual servers.
- Produces images to your specification.
- Doesn't do bare metal ☹
 - Look at
 - <https://github.com/google/netboot>
 - <https://github.com/jpetazzo/pxe>
 - <http://cobbler.github.io/>
 - <https://theforeman.org/>

Packer

- Very simple config file.
- Can use most provisioning systems.
 - Ansible, Puppet, Windows Shell, PowerShell, etc...
- Use the definition files you've already created.
- Gives you the true guarantee that no human has touched the image and everything done to the server is documented.



Step 3: Automate System Deployment

- You now have all the pieces for your server versioned and building.
- The next step is to automate their deployment and setting up things like networking.
- Most virtualisation has API's to do this.
- My preferred tool for this is Terraform (Done by the same company as packer)



Terraform

- Terraform is a tool for building, changing, and versioning infrastructure safely and efficiently.
- Terraform can manage existing and popular service providers as well as custom in-house solutions.

Terraform

- Description of infrastructure done via Text files.
- There's a pattern there 😊
- When you make a change it will work out the difference between what was done previously and current change.
- Can get it to return planned changes.



Terraform

- Can set up pretty much your entire infrastructure.
 - Networking
 - DNS
 - Servers
 - etc...

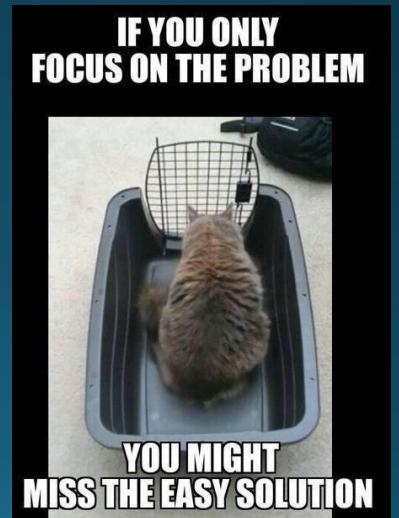


Terraform

- Outputs from one section can be used in a second.
 - i.e. IP's from servers it's setup can be used to create DNS entries later.
- This was one of the simplest new technologies I've learnt recently.
 - Had something up and working in about 2 hours. (Starting from no knowledge).

Step 4: Containerize Your Applications

- This area is still a work in progress.
- Docker seems to be the winner.
- More so in the Linux world than Windows.
 - Though you now have Windows native containers as well.
 - Microsoft working with Docker quite closely on this.



Container Advantages

- Immutability .
 - Images versioned.
 - Removes the “but it worked on my pc”.
 - No more difference between dev, testing, and prod.
 - Within reason i.e. DB
- Can be shared easily.
- Built off of text files (Dockerfiles).

Container Advantages

- Lower cost than virtualization.
- Base systems need fewer things installed which increases it's security.
- Allows splitting pieces up making it easier to update single piece.
- Simpler to test.
- Simpler roll back.
 - Within reason i.e. DB

What things aren't dealt with?

- Things that mutate:
 - Assets
 - Database
 - etc...
- Problems that already exist on the stack.



Data Not With Server

- Try to not keep the data with the server.
- Store images on NAS, GridFS, Ceph, etc...
 - When running in a cloud, use the provider's blob storage and similar storage.
- Relational DB's are a special case as you want them as close to the metal as possible.



My Current Pattern

My Current Pattern

- Image build with packer.
 - Docker installed and setup.
- Terraform for setup.
 - This depends on client and complexity.
- Docker Machine for key generation and Docker setup.
 - Allows controlling Docker remotely with shared SSL keys.
- Docker for all services running on servers.
 - Allows me to replicate entire stack locally when developing.

Some Interesting Things To Look At

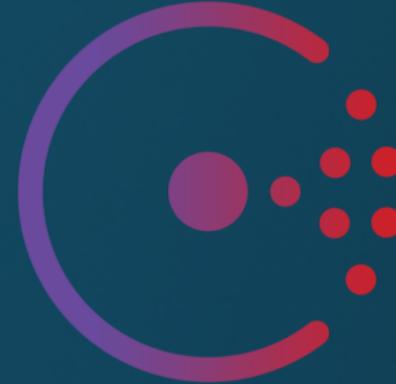


Chaos Monkey

- Developed by Netflix.
- Will randomly turn off services and servers.
- How sure are you about your failover and high availability.
- If you don't test you don't know.



Consul, etcd, etc.

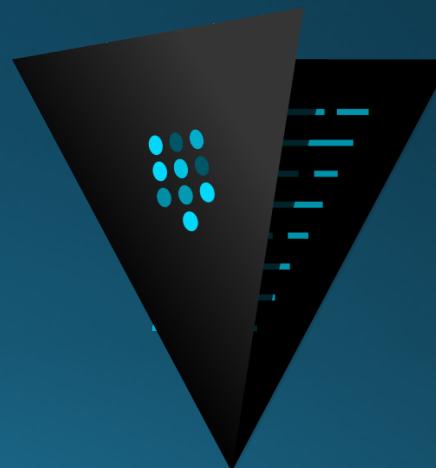


- Distributed fault tolerant DB.
- Aimed at storing config.
- Not built to be a database but as a way to store config and system state.
- Tooling to generate configs:
 - <https://github.com/hashicorp/consul-template>
 - <https://www.hashicorp.com/blog/introducing-consul-template/>



Vault

- Hashicorp, again...
- Centrally secure, store, and tightly control access to secrets across distributed infrastructure and applications.

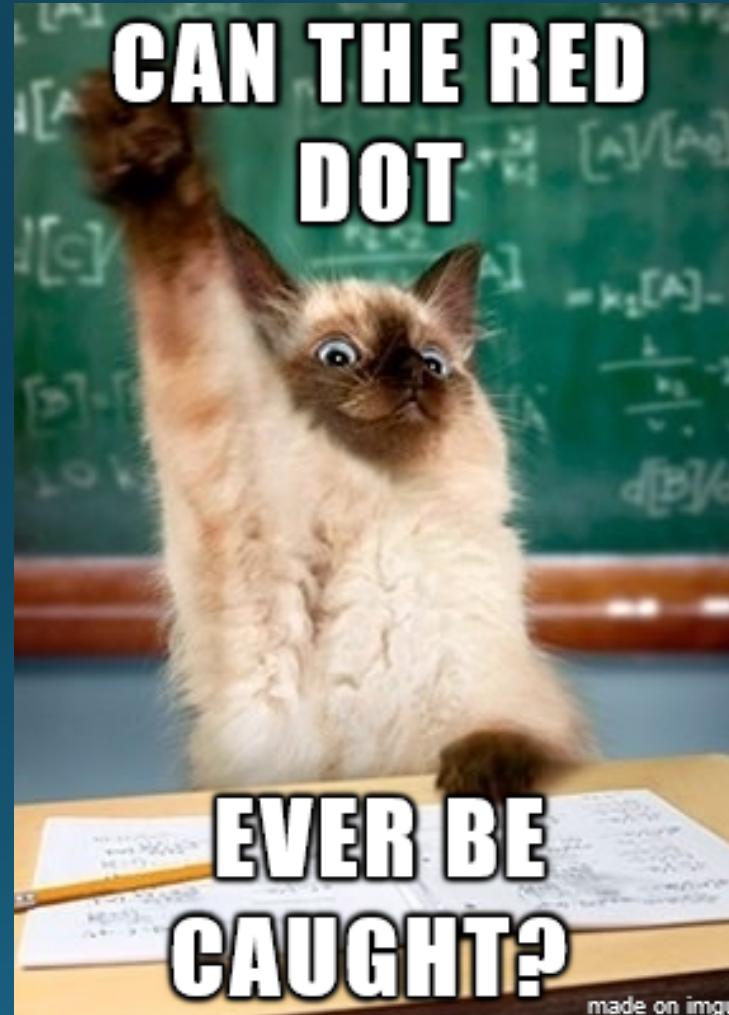


Rundeck



- Webased workflow and job scheduler.
- Can run things like Ansible etc...
- Centralized.
- Simple way for DevOps to provide GUI to Developers.
- <http://rundeck.org/>

Questions?



Thanks

